

# Python over Go

Kharitonskiy Vitaliy  
Prom.ua

# Images are important

- Eye candy
- Click-bait
- Small Details
- Text in images
- Marketing value
- Seeing is believing

# Images are not cheap

- Thousands of new images each day
- Storage/traffic worth \$\$\$
- Largest bandwidth consumer per page viewed, important for mobile site/app.
- Image processing is hard.



ReSRC.it



Kraken.io



Cloudinary

IMAGERESIZER




Sirv

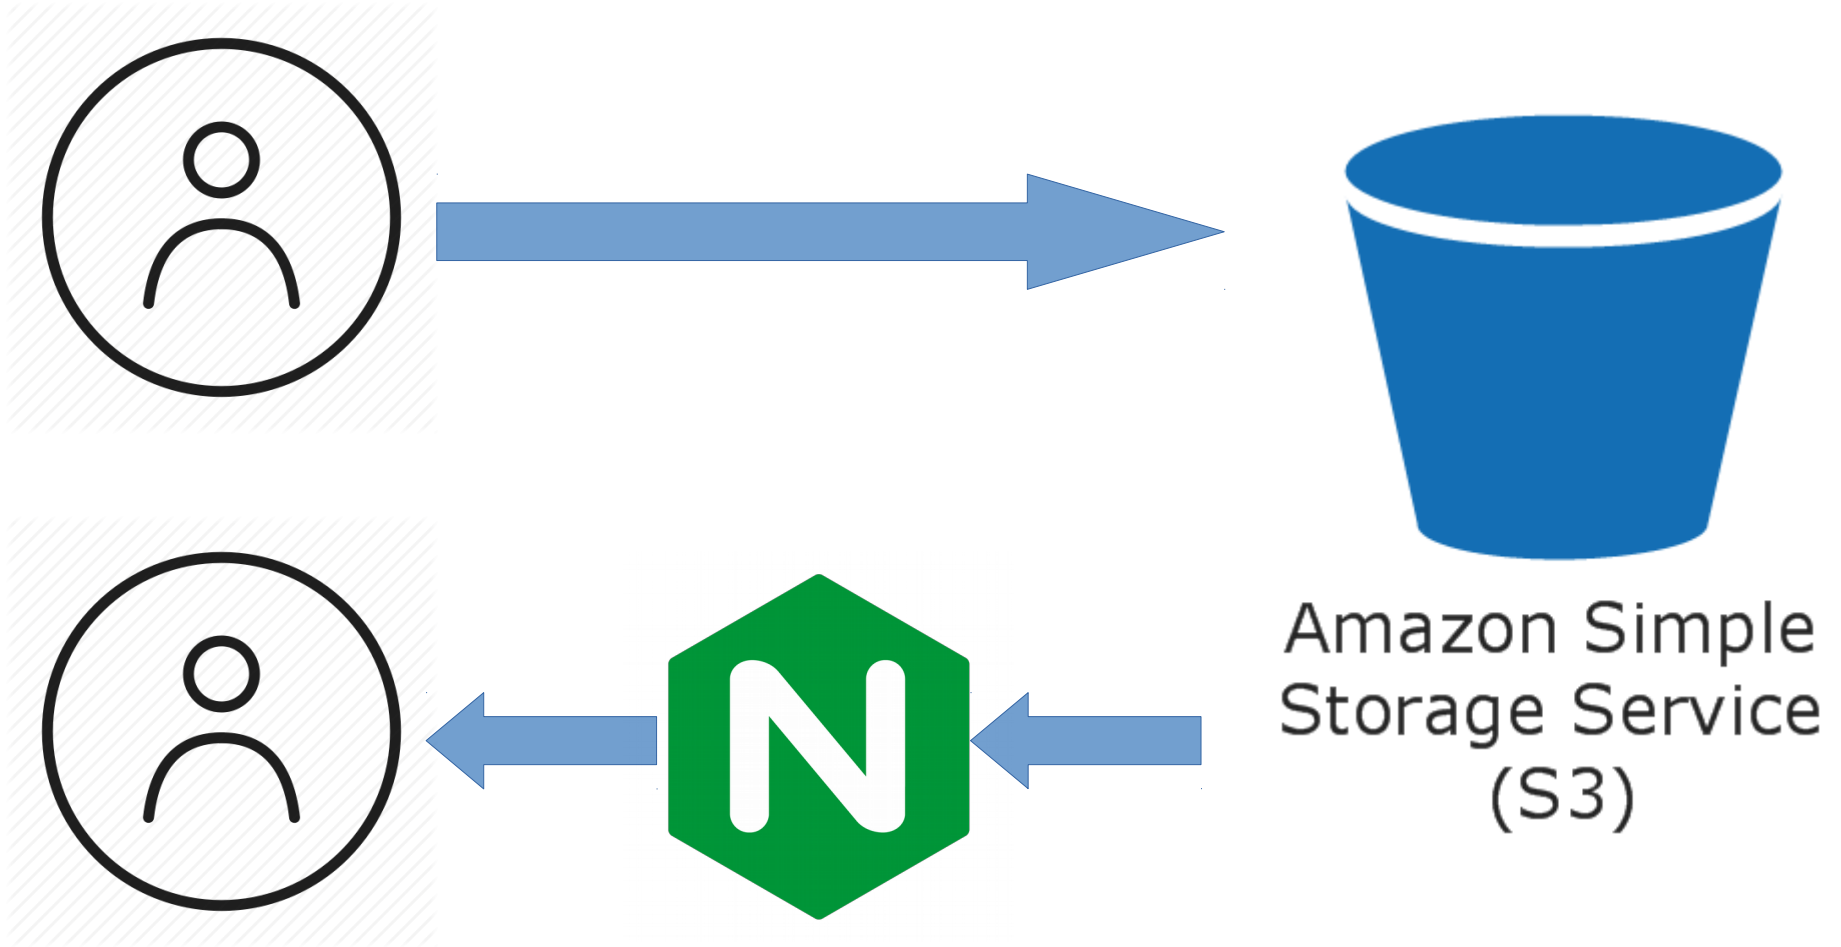
# Some numbers

- 6k images per second served at top load
- 120 Tb in Amazon
- 60 Tb in local cache
- 29 Riak nodes
- 6 nginx cache nodes
- 40 python processes across ~30 nodes for resize & watermark

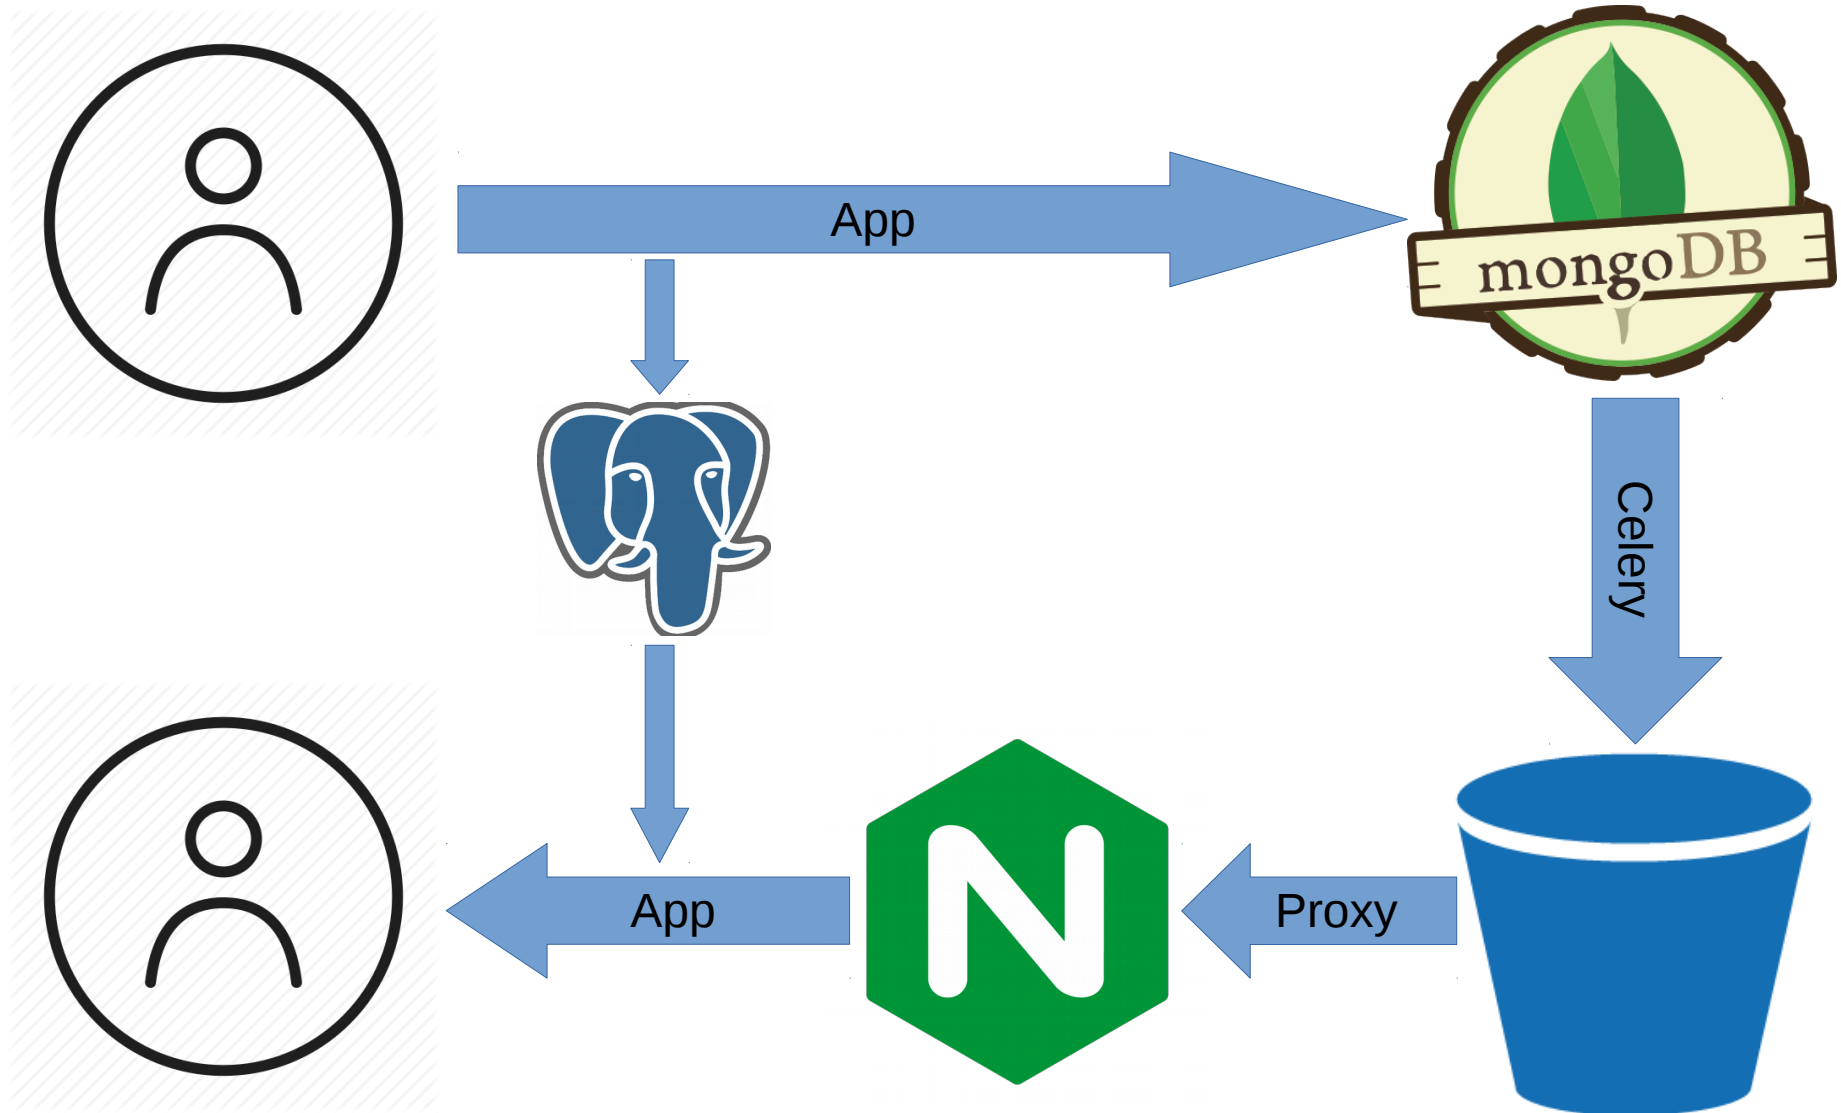
# Evolution of content delivery@Prom



# Evolution of content delivery@Prom




# Evolution of content delivery@Prom





# Evolution of content delivery@Prom



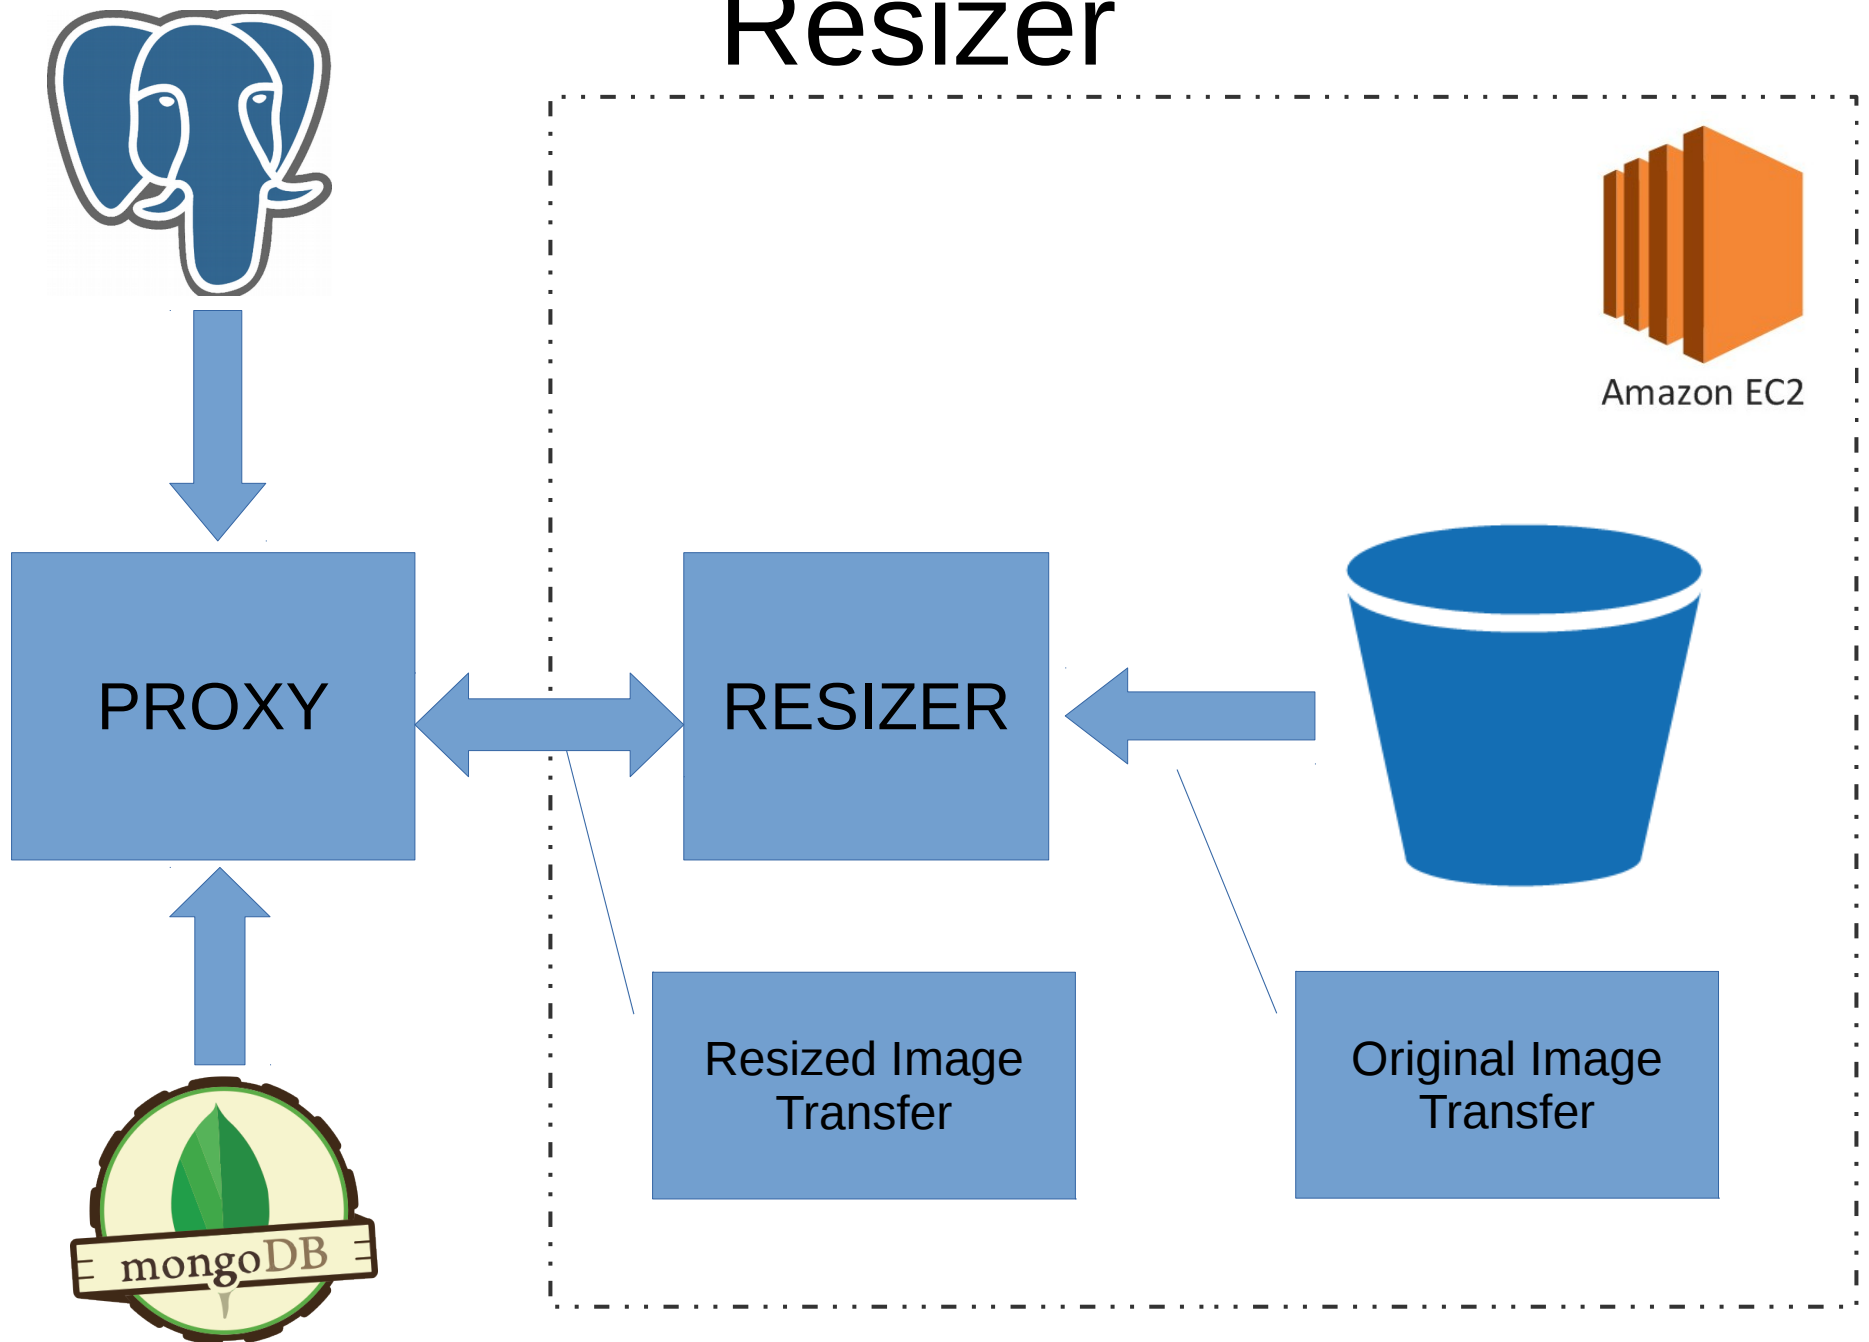
# Application Requirements

- Multiple data sources
- Unreliable local storage/frequent outages
- 80% IO / 20% CPU
- Reliable operations/deployment
- Shares logic with main app
- High load


# Why Golang?

- So hype, much fast, wow
- NOT Python
- Golang IS the new Python
- Because D'oh

# Resizer



# Resizer



# The Good

- Compiles to one binary, good for pre-container deployment
- HTTP is fast, Proxy part is stable
- Static compilation: no tests, zero downtime after deploy
- Gophers

# The Bad


- ImageMagick is not Pillow
- ImageMagick is slow
- Memory leaks in ImageMagick wrapper
- Weird concurrency/thread behavior
- Connection leaks everywhere
- Poor library/wrapper support(S3)
- DIY everywhere

# The Ugly

- return null, err
- if err == null
- if err == 'Not found'
- import "hell"
- Global Library Vars
- go func(){go func(){ go func yourself!!



# Simplified



# Microbenchmark(GO)

```
package main

import (
    "github.com/gographics/imagick"
    "io/ioutil"
)

func main() {
    blob, _ := ioutil.ReadFile("yuuuge.jpg")
    for i:=0; i<100; i++){
        mw := imagick.NewMagickWand()
        mw.ReadImageBlob(blob)
        mw.ResizeImage(200, 200, imagick.FILTER_LANCZOS2, 0.5)
        mw.Destroy()
    }
}
```

# Microbenchmark(Python)

```
from PIL import Image
from io import BytesIO

data = BytesIO(open('yuuuge.jpg', 'rb').read())

for i in range(100):
    with Image.open(data) as image:
        image.resize((200, 200)).save('result_resize.jpg', 'JPEG')
```

# Results

4712x3328 JPG resize to 200x200  
100 passes


**PYTHON**

28s total  
280 ms per pass  
3.5 resizes per  
second

**GOLANG**

30s total  
300 ms per pass  
3 resizes per  
second

# Results



# Why Python?

- AsyncIO for content proxy
- Pillow means compatibility with App
- Zero bus factor
- No more Mr. return err

# The Good

- PIL is fast
- PIL does not leak memory
- Simple, understandable concurrency behaviour
- Great libraries for everything(aiobotocore, aiopg, motor)
- Readable code

# The Bad

- Deployment is more complex(solved with containers)
- No static compilation(solved with tests/integration tests)
- CPU blocks IO(solved with more python processes)



# The Ugly



# The Real Ugly

- Request retry policy in botocore is configured via json file inside package folder

[https://github.com/boto/botocore/blob/develop/botocore/data/\\_retry.json](https://github.com/boto/botocore/blob/develop/botocore/data/_retry.json)

# The Real Ugly

- AsyncIO.sleep until loop is empty in case of delayed tasks inside handler

```
async def test_empty_image_fill(cli):
    resp = await cli.get('/7_w200_h200_test_empty.jpg')
    body = await resp.read()
    assert len(body) != 0
    assert resp.headers.get('X-Image-Source') == 'Resized original: local'
    # Wait until thumbnail is transferred to local cache
    await asyncio.sleep(0.4)
```

# CPU(BEFORE)

50000% = 50 cores = 8~9 servers at max load

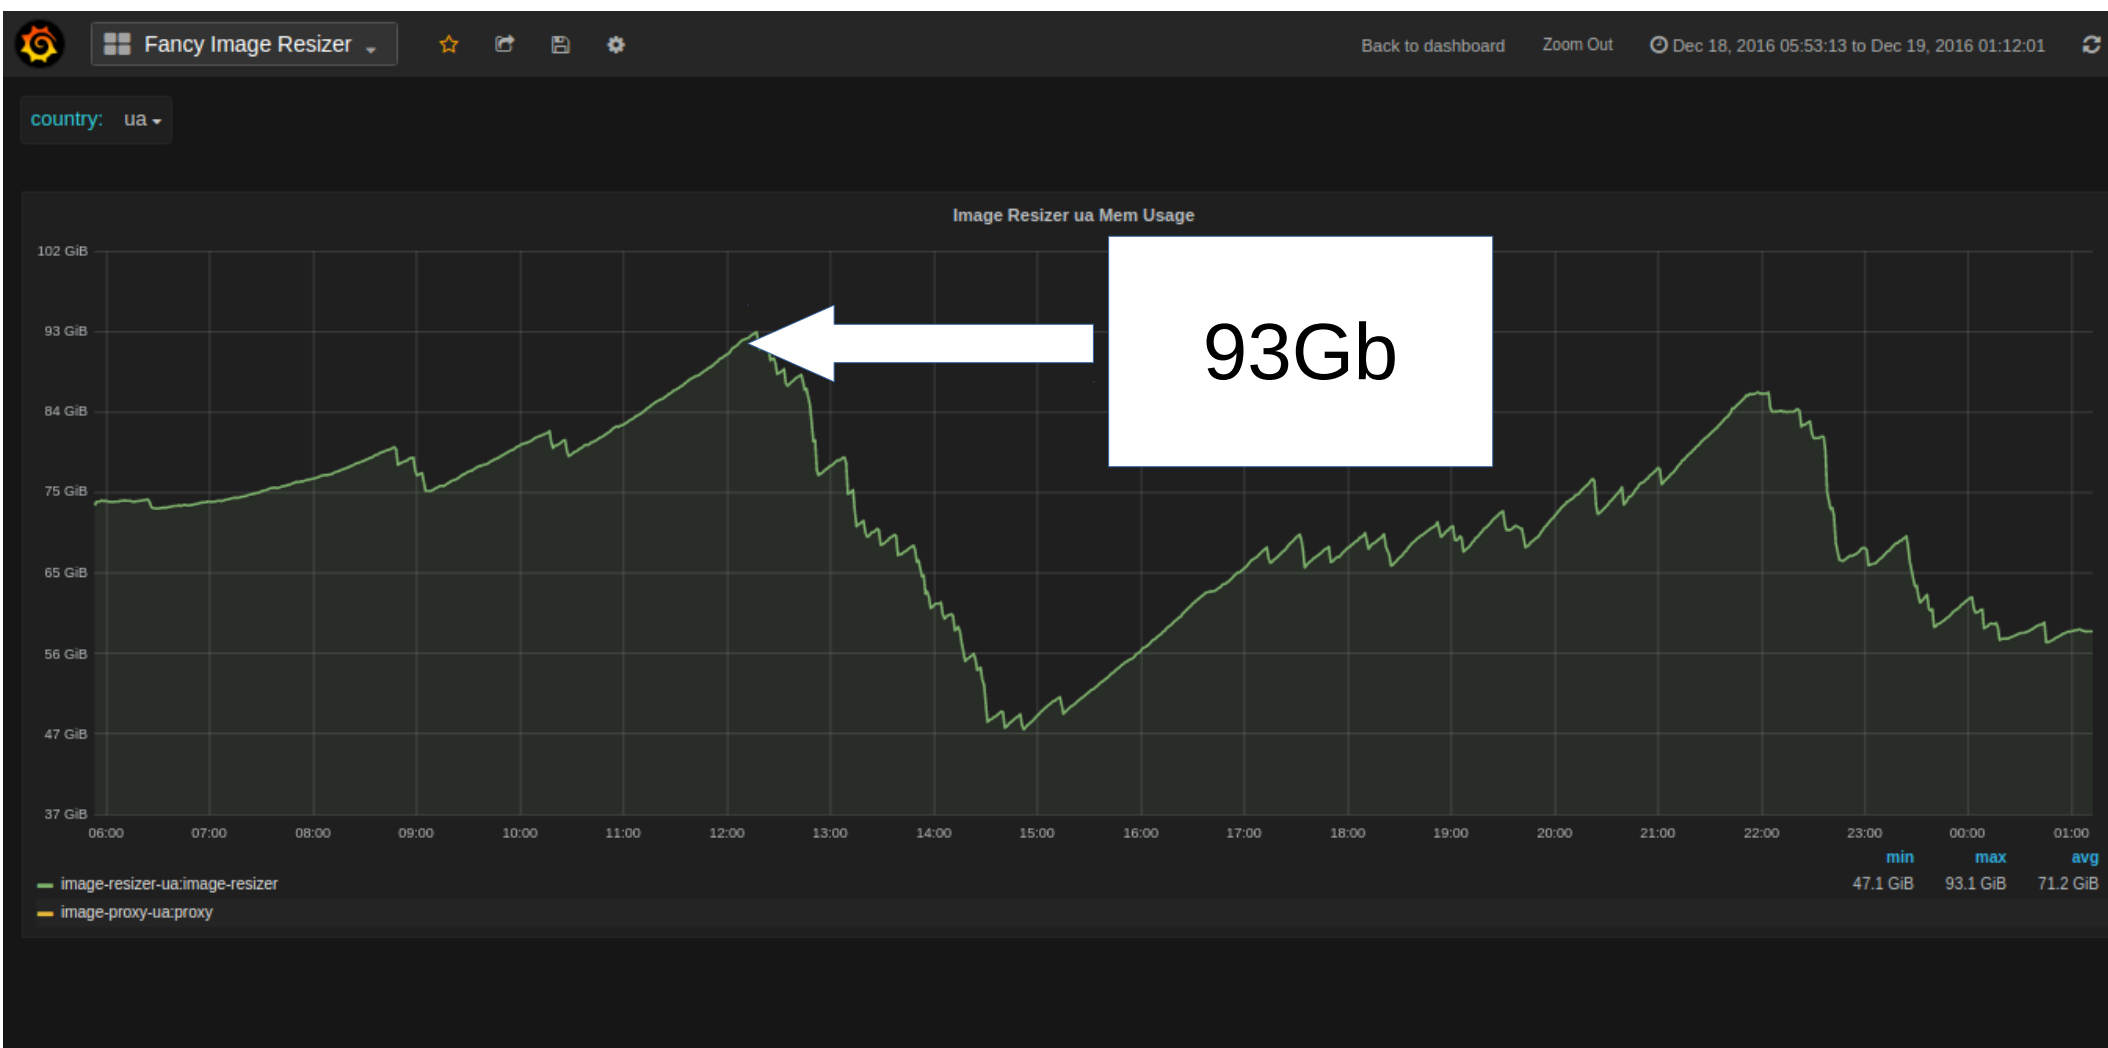


# CPU(AFTER)

1500% = 15 cores = 2~3 servers at max load



# Memory(BEFORE)



# Memory(AFTER)



# How we test

- pytest
- pytest-aiohttp
- GitLab CI
- Vagga - <https://github.com/tailhook/vagga>
- PostgreSQL
- MongoDB
- MinioGo
- ZERO Mocks



# How we deploy

- GitLab CI
- Vagga - <https://github.com/tailhook/vagga>
- Lithos - <https://github.com/tailhook/lithos>
- Verwalter - <https://github.com/tailhook/verwalter>

# How we measure

- Cantal - <https://github.com/tailhook/cantal>
- Carbon
- Graphite
- Graphana
- Logstash
- Kibana

# What is measured

- CPU
- Memory
- S3 Request time/count
- Riak Request time/count
- DB Request time/count
- Number of resizes
- Number of watermarks
- Number of hot/cold cache responses

# Measurement Rule of Thumb

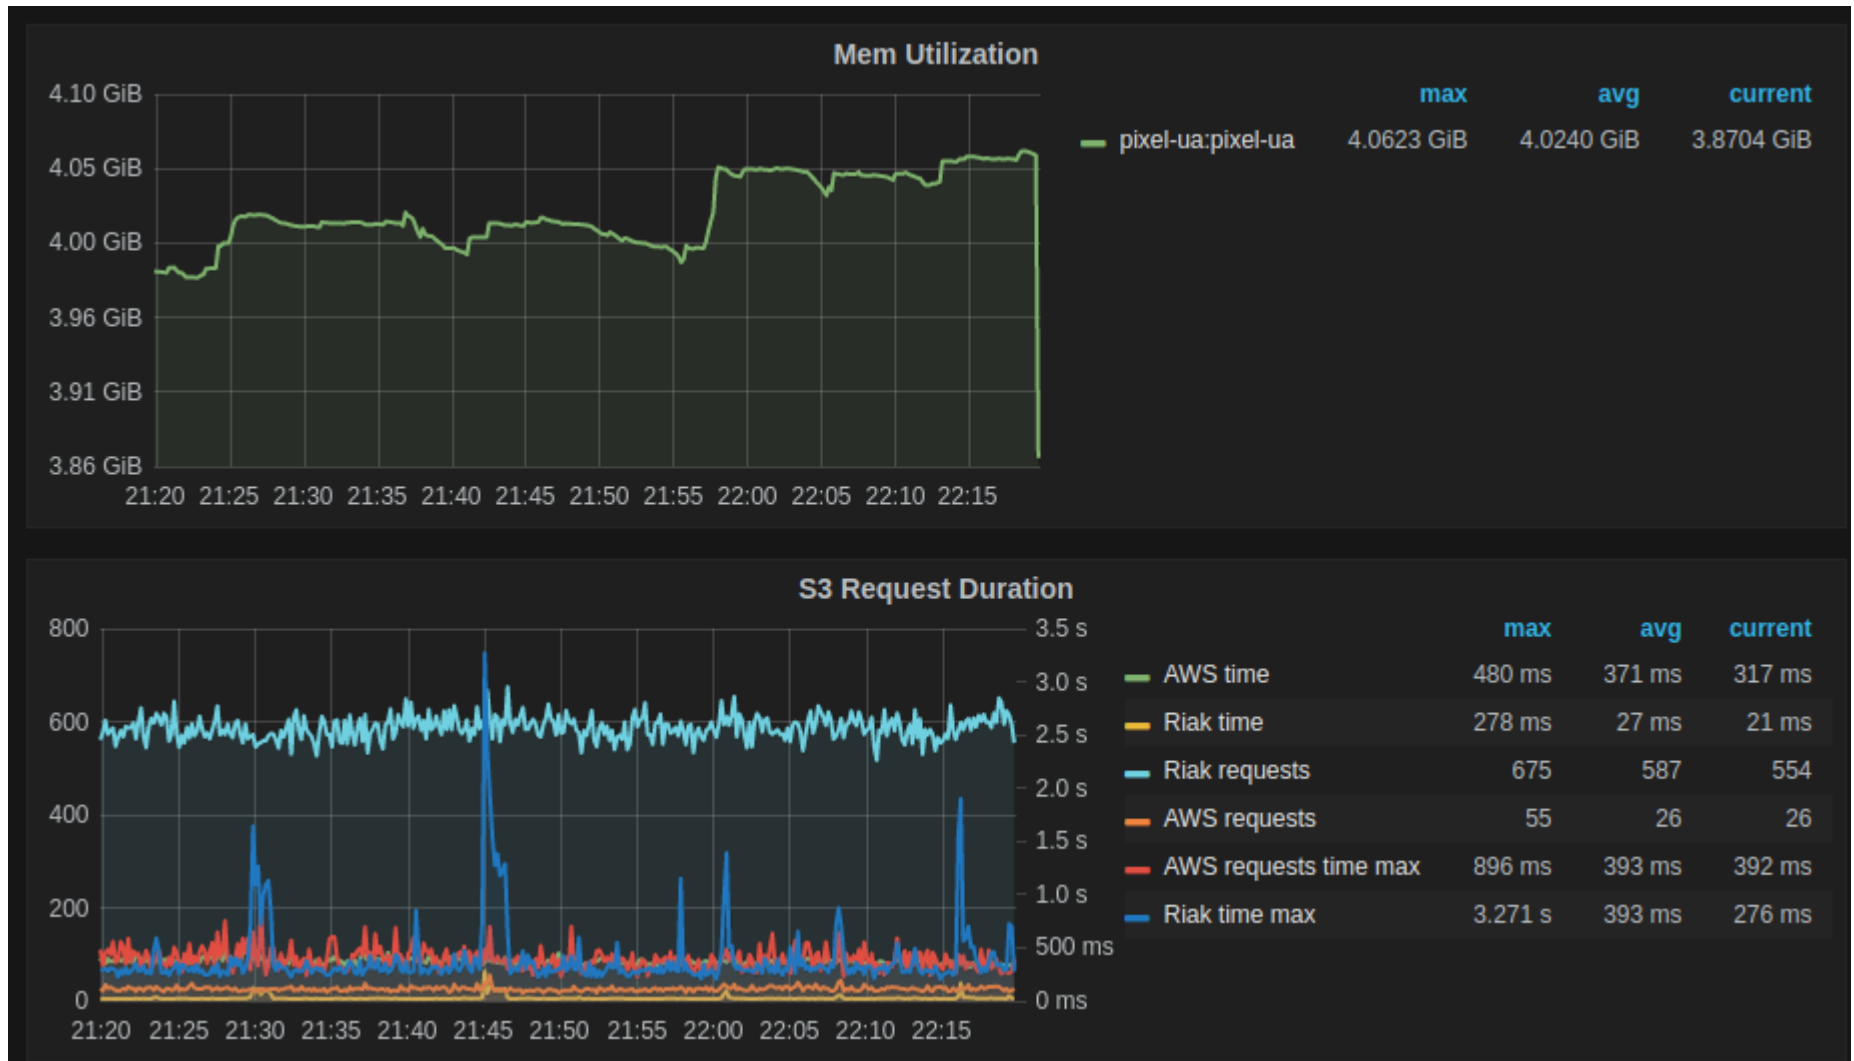
*What I do not see on the dashboard is probably not happening right now. Same applies to errors.*

*Confucius*


# Dashboard



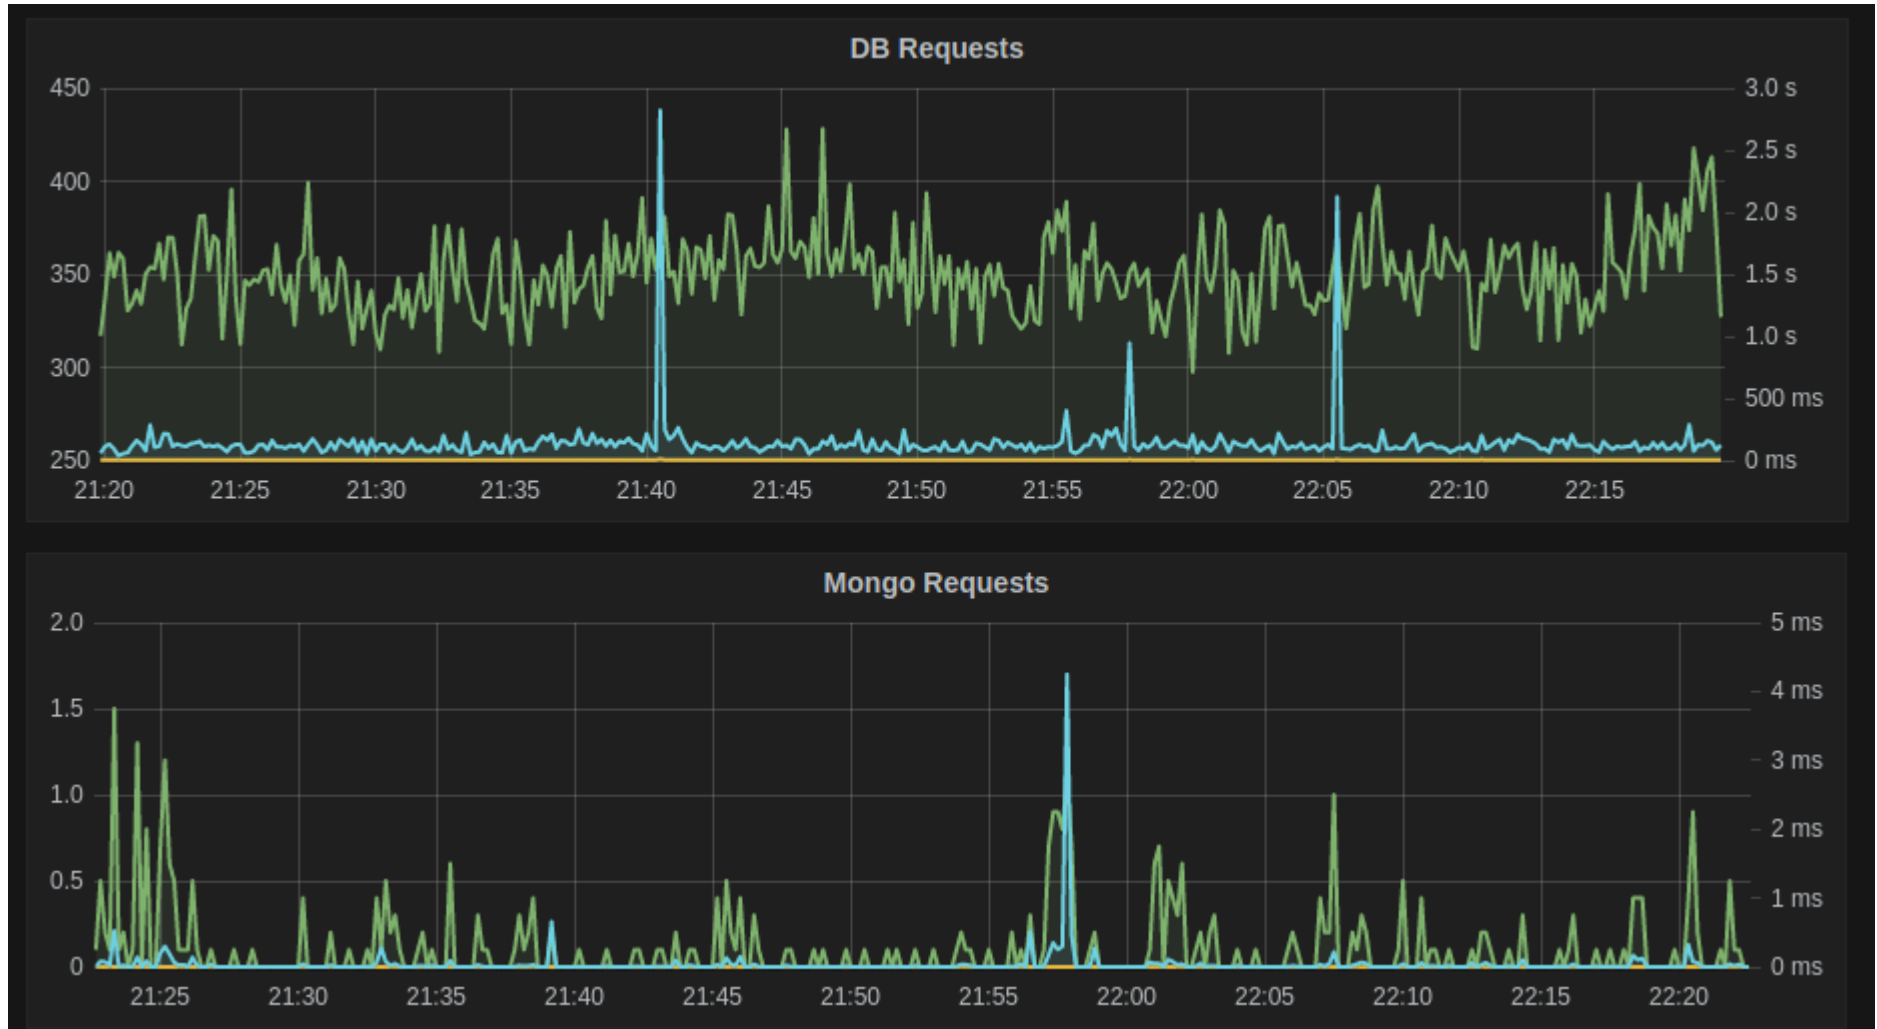
# Dashboard



# Dashboard




# Dashboard

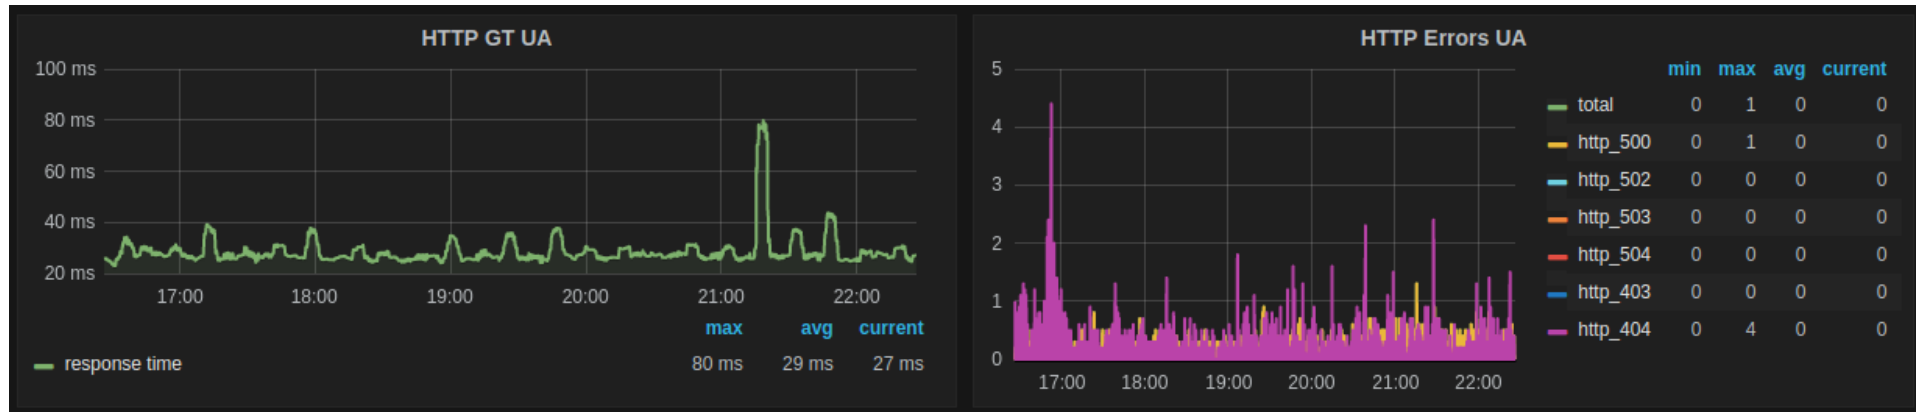





# Dashboard




# Dashboard



# Dashboard




# Dashboard



# Dashboard



# Dashboard



# Lessons learned

- Optimize the path for most popular size
- Read the code
- Go safe(aiopg vs asyncpg)
- Retries are evil
- Timeout every IO
- Inspect every IO operation for avg/max time
- DB connection for each request, use PgBouncer, use semaphore

# Tricks

- Do not touch the image
- Calculate format with mimetype, not with PIL
- Use Image.draft when possible
- Use Pillow SIMD with AXV2 compilation for speed
- Calculate source image quality before compressing the destination
- Nginx consistent hashing



# Pillow SIMD

- drop-in replacement for Pillow
- follows Pillow versions very closely
- can be compiled with AVX or AVX2
- very fast for resize, repack

<https://github.com/uploadcare/pillow-simd>

# Microbenchmark

```
from PIL import Image
from io import BytesIO

data = BytesIO(open('yuuuge.jpg', 'rb').read())

for i in range(100):
    with Image.open(data) as image:
        image.resize((200, 200)).save('result_resize.jpg', 'JPEG')
```

# Pillow SIMD vs Pillow

4712x3328 JPG resize to 200x200  
100 passes

**SIMD**

18s total  
180 ms per pass  
5.6 resizes per  
second

**VANILLA**

28s total  
280 ms per pass  
3.5 resizes per  
second

# Image.draft

- configure image reader to read resized thumbnail
- very fast as no Python code is involved
- keeps aspect ratio
- result size can be close to desired

# Microbenchmark

```
from PIL import Image
from io import BytesIO

data = BytesIO(open('yuuge.jpg', 'rb').read())

for i in range(100):
    with Image.open(data) as image:
        image.draft('RGB', (200, 200))
        image.save('result_draft.jpg', 'JPEG')
```

# Draft vs Resize

4712x3328 JPG resize to 200x200  
100 passes

**DRAFT**

12s total  
12 ms per pass  
83 resizes per  
second

**RESIZE**

28s total  
280 ms per pass  
3.5 resizes per  
second

# Image quality

- You have a 60 kb JPEG
- You open it with Pillow
- You save it to JPEG with quality=100
- What size will the result file have ?

**IT DEPENDS!**



# IT DEPENDS!

- Quality is not quality, it is a compression factor
- Given a JPEG with  $Q=75$ , compress it with  $Q=100$ , you get a bigger image
- Calculate quality with magic tables:  
<https://github.com/vharitonsky/imagequality>
- Compress image with `quality='keep'` if it is already compressed ( $Q \sim 60-75$ )

# Dyakuyu!

Embarrassing questions,  
embarrassing comments, libvips  
recommendations, etc.



# Dyakuyu!

Embarrassing questions,  
embarrassing comments, libvips  
recommendations, etc.

